



THE LINUX FOUNDATION
OPEN SOURCE SUMMIT
EUROPE

Demystifying memory: A Practical Tutorial on Managing & Optimizing memory in Zephyr

Marko Sagadin, IRNAS

About speaker

- Electrical (electronics) engineer by education
- Embedded systems engineer by profession
- PCB designer / assembly line engineer -> firmware engineer
- Consumer electronics, medical devices, IoT systems, TinyML...
- Secondary areas: tooling, developer environment, CI, processes, etc.






Your connected solution
development partner.

Maribor, Slovenia



Who is IRNAS?

- Cross-disciplinary team of 18 experts in connected product design and production
- 10 years of development of connected hardware products
- Full in-house infrastructure for innovating and manufacturing new products within days




Motivation

- An IoT device sending data via Wi-Fi
- PCB would contain nRF53 SOC, nRF7002 Wi-Fi companion IC and external flash chip.
- Only UI are an elnk display and two buttons
- On cloud side the Azure's IoT Hub service (application protocol built on top of MQTT)
 - Data is sent in strings, representing JSON objects
- Provisioning of Wi-Fi credentials will use BLE
- LVGL is used to render graphics
- External flash chip used to store assets and second image slot
- Code relocation and execute in place (XIP) were expected

Motivation

- Due to the number of features the higher memory usage was expected
- But how much?
- First good approximation was analysis of the *nrf/samples/lib/azure_iot_hub* sample.
- 61 KB FLASH and 69 KB RAM available
- Second MCUboot partition was already on external flash
- LTO was already enabled



```

MCUboot:
Memory region      Used Size  Region Size  %age Used
      FLASH:      46288 B      64 KB      70.63%
      RAM:        21328 B      48 KB      43.39%
      IDT_LIST:    0 GB      32 KB      0.00%

TF-M:
Memory region      Used Size  Region Size  %age Used
      FLASH:      129976 B    130560 B    99.55%
      RAM:        38996 B      48 KB      79.34%

Azure IoT Hub sample:
Memory region      Used Size  Region Size  %age Used
      FLASH:      702060 B     768 KB     89.27%
      RAM:        338376 B     400 KB     82.61%
      IDT_LIST:    0 GB      32 KB      0.00%

```


Flash and RAM

- Memory regions, not specific memory device technologies!
- FLASH
 - Read-only memory
 - Non-volatile – retains data after power loss
 - Stores program code (.text), read-only data (.rodata), and initial values for initialized data (.data)
- RAM
 - Read-write memory
 - Volatile – loses contents on power loss
 - Holds initialized data (.data), zero-initialized data (.bss), dynamic memory (heap), and call stacks (stack)

```
#include <stdint.h>
#include <stdio.h>
#include <stdlib.h>

/* Copied from FLASH to RAM at startup */
/* (Initialized global variable stored in .data section) */
uint8_t data_section = 42;

/* Stored in FLASH (read-only memory) */
/* (Constant global variable in .rodata section) */
const uint8_t rodata_section = 42;

/* Zero-initialized global variable stored in RAM */
/* (Uninitialized globals go to .bss section) */
uint8_t bss_section;

int main(int argc, char *argv[])
{
    /* Allocated on stack at runtime */
    /* (Local automatic variable inside main's stack frame) */
    uint8_t stack_var = 42;

    /* Allocated dynamically on heap at runtime */
    uint8_t *heap_section = malloc(sizeof(uint8_t));

    /* String literal stored in FLASH (.rodata section) */
    printf("Hello World\n");

    return 0;
}
```

Memory management in Zephyr

- `k_heap_*` API
 - Thread safe heap allocator
 - Timeout arguments
 - Built on top of `sys_heap_*`
- `k_mem_slab_*` API
 - Fixed-sized block allocator
- System Heap
 - heap allocator with libc-like API (`k_malloc`, `k_free`, etc.)
 - Wrapper around `k_heap`
 - It's size is set with `CONFIG_HEAP_MEM_POOL_SIZE`

```
/* zephyr/kernel/mempool.c, line 62 */  
K_HEAP_DEFINE(_system_heap, K_HEAP_MEM_POOL_SIZE);  
  
/* zephyr/subsys/net/ip/tcp.c, line 76 */  
K_MEM_SLAB_DEFINE_STATIC(tcp_conns_slab, sizeof(struct tcp),  
                           CONFIG_NET_MAX_CONTEXTS, 4);
```

Memory management in Zephyr

- Both macros are expanded to an array and a structure that references it.
- Prefixes: `kheap__*`, `_k_mem_slab_buf_*`
- The size of arrays is determined at compile time, and the reserved space **is included in the RAM usage report**.
- There are no “hidden” consumers, that are unaccounted for.
- Caveat: where `libc`’s `malloc` allocates from depends on the choice of `libc`.

```
/* zephyr/kernel/mempool.c, line 62 */
char kheap__system_heap[K_HEAP_MEM_POOL_SIZE];

struct k_heap__system_heap = {
    .heap = {
        .init_mem = kheap__system_heap,
        .init_bytes = K_HEAP_MEM_POOL_SIZE,
    },
};

/* zephyr/subsys/net/ip/tcp.c, line 76 */
size_t size = CONFIG_NET_MAX_CONTEXTS * sizeof(struct tcp);
static char _k_mem_slab_buf_tcp_conns_slab[size];

static struct k_mem_slab tcp_conns_slab = {
    .buffer = _k_mem_slab_buf_tcp_conns_slab,
    /* Other struct members... */
};
```

Very simplified code after pre-processor macro expansion

Useful tools and approaches for memory analysis

ROM and RAM report

- `west build -t rom_report`
- `west build -t ram_report`
- Sysbuild: `add -d build/<app dir>`

```
— wpa_supp_event_info      136    0.03% 0x000b0b98 rodata
— wpa_supp_ops.lto_priv.0  148    0.03% 0x2000ce4c datas
— wpa_supp_ready_sem       24    0.00% 0x2000d96c k_sem_area
— wpa_supp_status_work     48    0.01% 0x2000c4c0 datas
— wpa_supplicant_mutex     20    0.00% 0x2000d87c k_mutex_area
— wpas_api_ctrl            56    0.01% 0x20012c8c bss
— z_arm_tls_ptr             4    0.00% 0x20012e74 bss
— z_idle_stacks            320   0.06% 0x2003a810 noinit
— z_idle_threads          224   0.04% 0x2000de58 bss
— z_interrupt_stacks      2048  0.38% 0x20036038 noinit
— z_main_stack            6144  1.15% 0x2003a950 noinit
— z_main_thread           224   0.04% 0x20010058 bss
— z_malloc_heap            12    0.00% 0x20012e54 bss
— z_malloc_heap_mutex     20    0.00% 0x2000cb30 datas
— z_sys_post_kernel        1    0.00% 0x20015ade bss
— z_thread_monitor_lock    4    0.00% 0x2001256c bss
— zep_wifi_bh_q           256   0.05% 0x20012438 bss
— zep_wifi_intr_q         256   0.05% 0x20012338 bss
— zep_work_item           3200  0.60% 0x20013f58 bss
```

Binutils

- nm and addr2line
- They work most of the time

```
# Usecase: find type and location of a symbol by its name
arm-zephyr-eabi-nm -l build/<app>/zephyr/zephyr.elf | grep <symbol>

# Usecase: find location of a symbol by its address
arm-zephyr-eabi-addr2line -e build/<app>/zephyr/zephyr.elf <address>
```

Punccover

- Creates report in Web GUI
 - Code size
 - Variable size
 - Worst case statck analysis
- `west build -t puncover`
 - If using Sysbuild add `-d build/<app dir>`
 - `CONFIG_STACK_USAGE`
- What do the headers mean?
 - Remarks:
 - `x-module` – function is called from another file
 - `float` – function eventually calls a fp function
 - Code: `.text` section
 - Static: `.data`, `.rodata`, `.bss`

home / user / workdir / zephyr

Name	Remarks	Code	Static
arch		4,340	89
drivers		30,300	88,062
include		4,380	4
kernel		24,778	53,910
lib		18,946	3,408
modules		14,604	99,741
soc		594	68
subsys		118,402	58,296
Σ over all (8 folders, 0 files)		216,344	303,578

home / user / workdir / zephyr / kernel / mempool.c

Name	Remarks	Stack	Code	Static
k_aligned_alloc		8	108	
k_free	x-module	0	14	
k_malloc	x-module	0	8	
k_thread_system_pool_assign	x-module	0	12	
z_heap_aligned_alloc		32	104	
z_thread_aligned_alloc	x-module	16	48	
Σ 6 functions			294	
_system_heap				24
kheap_system_heap				40,144
Σ 2 variables				40,168
Σ over all (6 functions, 2 variables)			294	40,168

Thread analyzer

- Shows maximum stack usage for each thread at runtime
- Very versatile and easy to setup

```
CONFIG_THREAD_ANALYZER=y
CONFIG_THREAD_ANALYZER_USE_LOG=y
CONFIG_THREAD_ANALYZER_AUTO=y
CONFIG_THREAD_ANALYZER_AUTO_INTERVAL=5
CONFIG_THREAD_NAME=y
```

```
tcp_work      : STACK: unused 864 usage 160 / 1024 (15 %); CPU: 0 %
               : Total CPU cycles used: 0
shell_uart    : STACK: unused 5472 usage 672 / 6144 (10 %); CPU: 0 %
               : Total CPU cycles used: 10426
sysworkq      : STACK: unused 3864 usage 232 / 4096 (5 %); CPU: 0 %
               : Total CPU cycles used: 36
nrf70_intr_wq : STACK: unused 1448 usage 600 / 2048 (29 %); CPU: 0 %
               : Total CPU cycles used: 510
nrf70_bh_wq   : STACK: unused 824 usage 1224 / 2048 (59 %); CPU: 0 %
               : Total CPU cycles used: 90
logging       : STACK: unused 1544 usage 504 / 2048 (24 %); CPU: 0 %
               : Total CPU cycles used: 3790
idle          : STACK: unused 248 usage 72 / 320 (22 %); CPU: 99 %
               : Total CPU cycles used: 3265937
main          : STACK: unused 4792 usage 1352 / 6144 (22 %); CPU: 0 %
               : Total CPU cycles used: 1713
```

Heap and memory slab analysis

- System heap run time statistics
 - CONFIG_SYS_HEAP_RUNTIME_STATS
 - sys_heap_runtime_stats_get
 - Can also be used with k_heap
 - Shell command: kernel heap
- Memory slab's info member
 - num_blocks, block_size, num_use, max_used
 - CONFIG_MEM_SLAB_TRACE_MAX_UTILIZATION

```
#include <zephyr/kernel.h>
#include <zephyr/sys/sys_heap.h>

extern struct k_heap heap;
extern struct k_mem_slab slab;

void print_usage(void)
{
    k_spinlock_key_t key;
    struct sys_memory_stats stats;
    struct k_mem_slab_info info;

    key = k_spin_lock(&heap->lock);
    sys_heap_runtime_stats_get(&heap->heap, &stats);
    k_spin_unlock(&heap->lock, key);

    key = k_spin_lock(&slab->lock);
    info = slab.info;
    k_spin_unlock(&slab->lock, key);

    printk("Heap stats\n"
           "\t%zu free, %zu allocated, %zu max allocated",
           stats.free_bytes,
           stats.allocated_bytes,
           stats.max_allocated_bytes);

    printk("Slab stats\n"
           "%u blocks, %zu block size, %u used, %u max used",
           info.num_blocks,
           info.block_size,
           info.num_used,
           info.max_used);
}
```


Mbed TLS heap analysis

- Allocates from a static buffer.
- `CONFIG_MBEDTLS_HEAP_SIZE` controls size.
- Uses it's own memory allocator with it's own usage tracking logic.
- Depending on the type of Zephyr application you will need to configure different options.

```
# Zephyr project
CONFIG_MBEDTLS_SHELL=y
CONFIG_MBEDTLS_DEBUG=y
CONFIG_MBEDTLS_DEBUG_C=y
CONFIG_MBEDTLS_MEMORY_DEBUG=y


# For Zephyr project with TF-M also add:
CONFIG_TFM_CMAKE_BUILD_TYPE_RELWITHDEBINFO=y
# or
CONFIG_DEBUG=y

# NCS project with TF-M
# Compile ${ZEPHYR_BASE}/modules/mbedtls/shell.c
# in nrf/subsys/nrf_security/src/zephyr/CMakeLists.txt

# Shell
uart:~$ mbedtls heap
Maximum (peak): 52340 bytes, 111 blocks
Current: 0 bytes, 0 blocks
```

Network stack and nRF70 Wi-Fi

- Many “knobs” to turn, all affect memory consumption and throughput in different ways.
- RX and TX path are separately configurable
- net_pkt holds one or more net_bufs
- Driver for nRF7002 uses variable size.
- Interesting heaps to watch:
 - net_buf_mem_pool_rx_bufs
 - net_buf_mem_pool_tx_bufs
 - wifi_drv_ctrl_mem_pool
 - wifi_drv_dat_mem_pool



Optimization strategy

- Locate big spenders with the static tools.
 - Determine symbols that enable or tune them.
 - Determine if you even need them.
 - Tackle them first.
- Many subsystem have constrained variants
- Decide what you can miss and what is mandatory to have.
- Define what are your device workflows and possible code paths.
- Determining correct sizes of threads, heaps and mem slabs:
 - 1st estimation: Run your code through all possible code paths and then check report tools for maximum values.
 - 2nd estimation: keep device running for several days.
 - 3rd estimation: run several devices in the field and have them report back statistics.

Memory analysis demo

<https://github.com/MarkoSagadin/demistifying-memory-project>

Azure IoT Hub memory results

Test conditions:

- Sending 1kB of data (hard-coded string) 10 times every 20 sec,
- Received 1kB data sporadically, CJSON parses it
- Disconnection (sudden removal of the access point)
- No FOTA, no direct methods, little CJSON use.

Build variant	Used Flash	Max Flash	Free Flash	Used RAM	Max RAM	Free RAM
Original	707 KB	768 KB	61 KB	331 KB	400 KB	69 KB
Optimized 1	684 KB	768 KB	84 KB	226 KB	400 KB	174 KB
Optimized 2	668 KB	768 KB	100 KB	226 KB	408 KB	182 KB

What I learned?

- How to find things in the Zephyr.
- General rule: FLASH usage is “hard” feature bound , but RAM is “soft” feature bound (more tunable).
- External flash and execute in place (XIP) are essentially required for nRF7002 (STA applications). Nordic documents well how this can be done.
- If you can, keep heaps separate. Although they can waste RAM space, they make it much easier to reason about the system.

IRNAS's resources

- Found at: <https://github.com/irnas>
- To speed up our own development we made many different project templates, reusable scripts, documentation. They are public and free to use. Search for “irnas-” prefix.
- Contents:
 - Docker images
 - CI infrastructure
 - Zephyr project template
 - Custom East tool
 - Project and Development guidelines

